

scaler-clustering

November 29, 2024

0.1 About Scaler

Scaler is an online tech-versity offering intensive computer science & Data Science courses through live classes delivered by tech leaders and subject matter experts. The meticulously structured program enhances the skills of software professionals by offering a modern curriculum with exposure to the latest technologies. It is a product by InterviewBit.

0.1.1 Business Problem

To cluster a segment of learners from the Scaler database based on their job profile, company, and other relevant features, with the goal of profiling the best companies and job positions to work for. The resulting clusters should group learners with similar characteristics.

0.1.2 Dataset

Column Name	Description
Unnamed 0	Index of the dataset
Email_hash	Anonymised Personal Identifiable Information (PII)
Company_hash	This represents an anonymized identifier for the company, which is the current employer of the learner.
orgyear	Employment start date
CTC	Current CTC
Job_position	Job profile in the company
CTC_updated_year	Year in which CTC got updated (Yearly increments, Promotions)

Importing Required Libraries

```
[1]: import pandas as pd
import numpy as np
import re

from scipy import stats
from sklearn.impute import KNNImputer
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
```

```

from scipy.cluster.hierarchy import dendrogram, linkage, fcluster

from matplotlib import pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')

```

```
[2]: palette = ['#0000ff', '#ffffff', '#000000', '#99c2ff']
```

Read Dataset

```
[3]: df = pd.read_csv(r'scaler_clustering.csv', index_col=0)
df.head()
```

```
[3]:
```

	company_hash \		email_hash	orgyear	ctc \
0	atrgxnnt xzaxv		6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...	2016.0	1100000
1	qtrxvzwt xzegwgbg rxbxnta		b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10...	2018.0	449999
2	ojzwnvwnxw vx		4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9...	2015.0	2000000
3	ngpgutaxv		effdede7a2e7c2af664c8a31d9346385016128d66bbc58...	2017.0	700000
4	qxen sqghu		6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...	2017.0	1400000

	job_position	ctc_updated_year
0	Other	2020.0
1	FullStack Engineer	2019.0
2	Backend Engineer	2020.0
3	Backend Engineer	2019.0
4	FullStack Engineer	2019.0

```
[4]: print("Shape of the data: ", df.shape)
print("The Given Dataset has {} rows and {} columns".format(df.shape[0], df.
↪shape[1]))
print("Columns: ", df.columns.to_list())
```

Shape of the data: (205843, 6)

The Given Dataset has 205843 rows and 6 columns

Columns: ['company_hash', 'email_hash', 'orgyear', 'ctc', 'job_position', 'ctc_updated_year']

0.1.3 Shape

- The dataset comprises 205843 rows and 6 columns, representing a volume of data.
- Each row corresponds to information about the learners details.

0.1.4 Data Structure

```
[5]: df.describe()
```

```
[5]:
```

	orgyear	ctc	ctc_updated_year
count	205757.000000	2.058430e+05	205843.000000
mean	2014.882750	2.271685e+06	2019.628231
std	63.571115	1.180091e+07	1.325104
min	0.000000	2.000000e+00	2015.000000
25%	2013.000000	5.300000e+05	2019.000000
50%	2016.000000	9.500000e+05	2020.000000
75%	2018.000000	1.700000e+06	2021.000000
max	20165.000000	1.000150e+09	2021.000000

```
[6]: df.describe(include='object')
```

```
[6]:
```

	company_hash \
count	205799
unique	37299
top	nvnv wgzohrnrvzwj otqcxwto
freq	8337

	email_hash	job_position
count	205843	153279
unique	153443	1016
top	bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...	Backend Engineer
freq	10	43554

```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 205843 entries, 0 to 206922
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   company_hash           205799 non-null    object
1   email_hash             205843 non-null    object
2   orgyear                205757 non-null    float64
3   ctc                    205843 non-null    int64
4   job_position           153279 non-null    object
5   ctc_updated_year       205843 non-null    float64
dtypes: float64(2), int64(1), object(3)
```

memory usage: 11.0+ MB

```
[8]: df.isnull().sum()
```

```
[8]: company_hash      44
     email_hash        0
     orgyear           86
     ctc               0
     job_position     52564
     ctc_updated_year  0
     dtype: int64
```

```
[9]: print(df.duplicated().sum())
     df.drop_duplicates(inplace=True)
```

34

0.1.5 Dataset Information:

- **Data Consistency:** job_position column has more number of missing values, columns like orgyear, company_hash also has missing values.
- **Data Types:** Columns are classified into integer, float and object types.
- **Data Duplicates:** There are 132 rows in the record are duplicated.

Initial Cleanup

```
[10]: df['job_position'] = df['job_position'].apply(lambda x: re.sub('[^A-Za-z0-9_ ]+', '', x.lower()) if isinstance(x, str) else x)
     df['company_hash'] = df['company_hash'].apply(lambda x: re.sub('[^A-Za-z0-9_ ]+', '', x.lower()) if isinstance(x, str) else x)
     df['email_hash'] = df['email_hash'].apply(lambda x: re.sub('[^A-Za-z0-9 ]+', '', x.lower()) if isinstance(x, str) else x)
```

```
[11]: df['orgyear'] = df.groupby(['company_hash'])['orgyear'].transform(lambda x: x.
     ↪ffill().bfill())
     print(df['orgyear'].isnull().sum())
     df['orgyear'] = df['orgyear'].fillna(df['orgyear'].mode()[0])
```

70

```
[12]: # Convert the data types
     df['job_position'] = df['job_position'].astype('category')
     df['company_hash'] = df['company_hash'].astype('category')
     df['email_hash'] = df['email_hash'].astype('category')

     df['ctc'] = df['ctc'].astype('int')
     df['ctc_updated_year'] = df['ctc_updated_year'].astype('int')
     df['orgyear'] = df['orgyear'].astype('int')
```

0.1.6 Data Observation

```
[13]: df.shape
```

```
[13]: (205809, 6)
```

```
[14]: ##### Frequency of Email Id and Company Hash
# df.groupby(['email_hash', 'company_hash']).size().reset_index(name='count').
#   ↪sort_values('count', ascending=False).head()
df[['email_hash', 'company_hash']].value_counts().reset_index(name='count').
#   ↪sort_values('count', ascending=False).head()
```

```
[14]:                                     email_hash \
0  bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...
2  298528ce3160cc761e4dc37a07337ee2e0589df251d736...
3  6842660273f70e9aa239026ba33bfe82275d6ab0d20124...
1  3e5e49daa5527a6d5a33599b238bf9bf31e85b9efa9a94...
4  c0eb129061675da412b0deb15871dd06ef0d7cd86eb5f7...
```

	company_hash	count
0	oxej ntwyzgrgsxto rxbxnta	10
2	cvrhtbgbtznhb	9
3	ihvrwgbb	9
1	wgcxvb ntwyzgrgsxto	9
4	nyt a t oyvf sqghu	8

```
[15]: df[df['email_hash'].str.
#   ↪contains('bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7')]
```

```
[15]:                                     company_hash \
24129  oxej ntwyzgrgsxto rxbxnta
46038  oxej ntwyzgrgsxto rxbxnta
72415  oxej ntwyzgrgsxto rxbxnta
103145 oxej ntwyzgrgsxto rxbxnta
118076 oxej ntwyzgrgsxto rxbxnta
121825 oxej ntwyzgrgsxto rxbxnta
124840 oxej ntwyzgrgsxto rxbxnta
145021 oxej ntwyzgrgsxto rxbxnta
153402 oxej ntwyzgrgsxto rxbxnta
160472 oxej ntwyzgrgsxto rxbxnta
```

	email_hash	orgyear	ctc	\
24129	bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018	720000	
46038	bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018	720000	
72415	bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018	720000	
103145	bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018	720000	

118076	bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018	720000
121825	bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018	660000
124840	bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018	660000
145021	bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018	660000
153402	bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018	660000
160472	bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018	660000

	job_position	ctc_updated_year
24129	NaN	2020
46038	support engineer	2020
72415	other	2020
103145	fullstack engineer	2020
118076	data analyst	2020
121825	other	2019
124840	support engineer	2019
145021	fullstack engineer	2019
153402	devops engineer	2019
160472	NaN	2019

```
[16]: df.
      ↪sort_values(by=['email_hash', 'company_hash', 'orgyear', 'ctc', 'ctc_updated_year', 'job_position'],
      ↪ascending=True, inplace=True)
df['job_position'] = df.groupby(['email_hash', 'company_hash', 'ctc',
      ↪'ctc_updated_year'])['job_position'].transform(lambda x: x.ffill().bfill())
```

```
[17]: print(df.duplicated().sum())
df.drop_duplicates(inplace=True)
```

27308

0.1.7 Duplicate value check:

- **Frequency of Company/Email Hash:** It is observed that the for a individual learner there are multiple entires in the dataset with different job position name.
- **Data Types:** Columns are classified into integer, float and object types. DateTime columns are stored as Objects types.

```
[18]: print("No. of Unique Job Positions: ", df['job_position'].nunique())
print("No. of Null in Job Positions: ",df['job_position'].isnull().sum())
```

No. of Unique Job Positions: 931
No. of Null in Job Positions: 25332

```
[19]: ## Manual clean up of job positions

# jobs_df = df['job_position'].value_counts().to_frame()
# jobs_df.reset_index(inplace=True)
# jobs_df.columns = ['job_position', 'count']
```

```
# jobs_df.to_csv('job_positions.csv', index=False)

jobs_df = pd.read_csv('job_positions.csv')
jobs_df.head()
```

```
[19]:
```

	job_position	count	filler
0	backend engineer	43551	Backend Engineer
1	fullstack engineer	25975	Fullstack Engineer
2	other	18070	Others
3	frontend engineer	10417	Frontend Engineer
4	engineering leadership	6870	Technical lead

```
[20]: job_position_dict = dict(zip(jobs_df['job_position'], jobs_df['filler']))
# job_position_dict
```

```
[21]: df['job_position'] = df['job_position'].map(job_position_dict)
```

```
[22]: print("Null count: ", df['job_position'].isnull().sum())
print("Number of Uniques: ", df['job_position'].nunique())
```

```
Null count: 25335
Number of Uniques: 41
```

```
[23]: # Encode the categorical column
# Temporarily fill NaN values with a placeholder
df['job_position'].fillna('missing', inplace=True)

# Encode the categorical column
le = LabelEncoder()
df['job_position_encoded'] = le.fit_transform(df['job_position'])

# Replace the placeholder back to NaN
df['job_position_encoded'].replace(np.where(le.classes_ == 'missing')[0][0], np.
    ↳nan, inplace=True)

# Apply KNN imputer
imputer = KNNImputer(n_neighbors=5)
df['job_position_encoded'] = imputer.fit_transform(df[['job_position_encoded']])

# Decode the imputed values back to original categories
df['job_position'] = le.inverse_transform(df['job_position_encoded'].
    ↳astype(int))

# Drop the encoded column
df.drop(columns=['job_position_encoded'], inplace=True)
```

```
[24]: df.isnull().sum()
```

```
[24]: company_hash      39
      email_hash        0
      orgyear           0
      ctc               0
      job_position      0
      ctc_updated_year  0
      dtype: int64
```

```
[25]: print(df['company_hash'].nunique())
      df.dropna(subset=['company_hash'], inplace=True)
```

37299

```
[26]: df.isnull().sum()
```

```
[26]: company_hash      0
      email_hash        0
      orgyear           0
      ctc               0
      job_position      0
      ctc_updated_year  0
      dtype: int64
```

0.1.8 Exploratory Data Analysis (EDA)

```
[27]: # CTC Column checks
      print(df['ctc'].min(), df['ctc'].max())

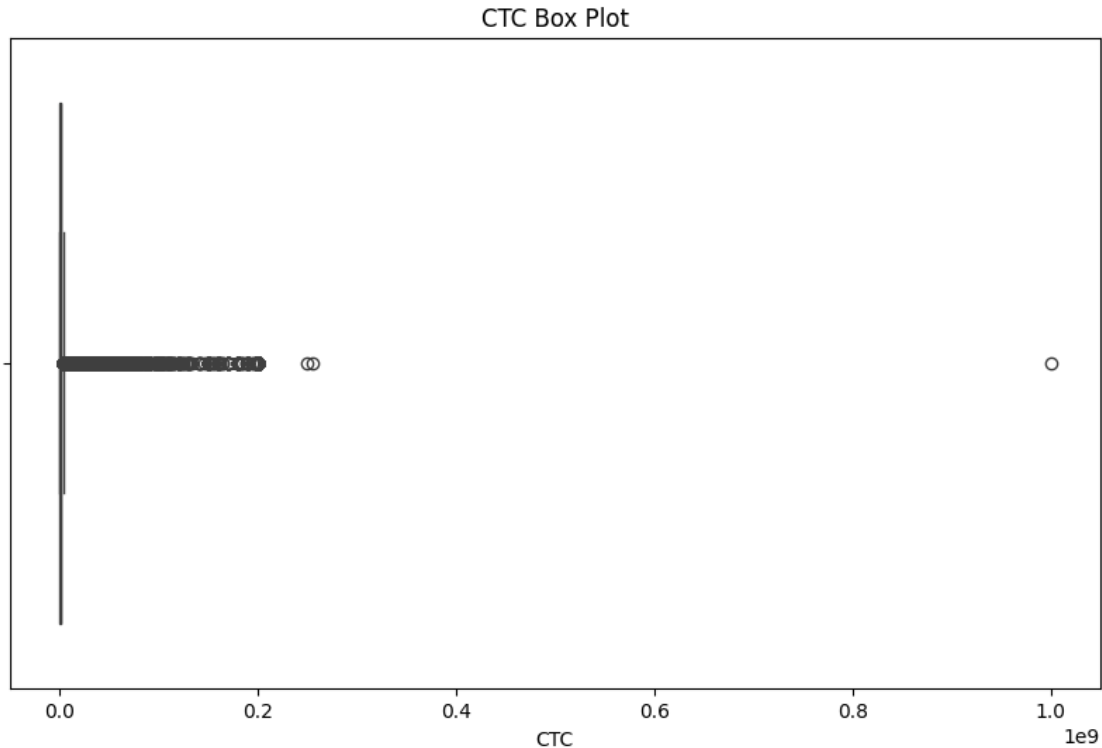
      # Calculate Percentiles
      lower_bound = df['ctc'].quantile(0.01)
      upper_bound = df['ctc'].quantile(0.99)

      print(f"Lower Bound: {lower_bound}, Upper Bound: {upper_bound}")
```

2 1000150000

Lower Bound: 35000.0, Upper Bound: 16000000.0

```
[28]: plt.figure(figsize=(10, 6))
      sns.boxplot(x=df['ctc'])
      plt.title('CTC Box Plot')
      plt.xlabel('CTC')
      plt.show()
```

```
[29]: df['CTC_zscore'] = stats.zscore(df['ctc'])

# Identify outliers (e.g., Z-score > 3 or < -3)
outliers = df[(df['CTC_zscore'] > 3) | (df['CTC_zscore'] < -3)]

print("Outliers based on Z-score:")
print("Number of outliers: ", outliers.shape[0])
print("Percentage of outliers: ", outliers.shape[0] / df.shape[0] * 100)
```

```
Outliers based on Z-score:
Number of outliers: 1499
Percentage of outliers: 0.8399547242550235
```

```
[30]: outliers.sort_values(by='CTC_zscore', ascending=True).head()
```

```
[30]:
```

	company_hash	email_hash \
168918	ftirro evqsg	55c7eef700e87fcd506c731be6dbbdc79bf709678c2eb...
50457	wgzwtznqxd	c240203d659ee5ef1a7ed9d8368ec86b7fc1d21e10701b...
15755	wgszxkvzn	405acce4415219a5001f40c37ca2e5f07d433a8e769641...
77447	wgszxkvzn	405acce4415219a5001f40c37ca2e5f07d433a8e769641...
141547	zgn vuurxwvmrt	a7b47e958b5a48f375cb74e23537b601396356545a19c5...

orgyear	ctc	job_position	ctc_updated_year	CTC_zscore
---------	-----	--------------	------------------	------------

168918	2002	39600000	Technical lead	2019	3.006501
50457	2017	39800000	Others	2020	3.022664
15755	2017	39900000	QA Engineer	2020	3.030746
77447	2017	39900000	Support Engineer	2020	3.030746
141547	2021	40000000	Frontend Engineer	2019	3.038828

```
[31]: outliers.sort_values(by='CTC_zscore', ascending=False).head()
```

```
[31]:
```

	company_hash \		email_hash	orgyear \
72925	whmxw rgxwo uqxcvnt rxbxnta			
117948		obvqnuqxwdwg		
3301	aveegaxr xzntqzvnxgzvr hzxctqoxnj			
9220	vayxxuv ntwyzgrgsxto ucn rna			
107466	ytfrtnn uvwpvqa tzntquqxot			

	ctc	job_position	ctc_updated_year	CTC_zscore
72925	1000150000	Frontend Engineer	2020	80.635227
117948	255555555	Frontend Engineer	2016	20.459371
3301	250000000	Frontend Engineer	2020	20.010388
9220	200000000	Co founder	2020	15.969541
107466	200000000	Data Analyst	2020	15.969541

```
[32]: df = df[df['ctc']>=1000000]
df = df[df['ctc']<=15000000]
```

```
[33]: print("Median: ",df['ctc'].median())
print("Mean: ",df['ctc'].mean())
```

```
Median: 1000000.0
Mean: 1370185.5501391143
```

```
[34]: # Create ctc into 3 categories
# df['ctc_category'] = pd.qcut(df['ctc'], q=3, labels=['low', 'medium', 'high'])
df['ctc_category'] = pd.cut(df['ctc'], bins=[0, 1000000, 5000000, 15000000],
↳labels=['low', 'medium', 'high'])
```

```
[35]: df['ctc_category'].value_counts()
```

```
[35]: ctc_category
low      91187
medium   78056
```

```
high      3277
Name: count, dtype: int64
```

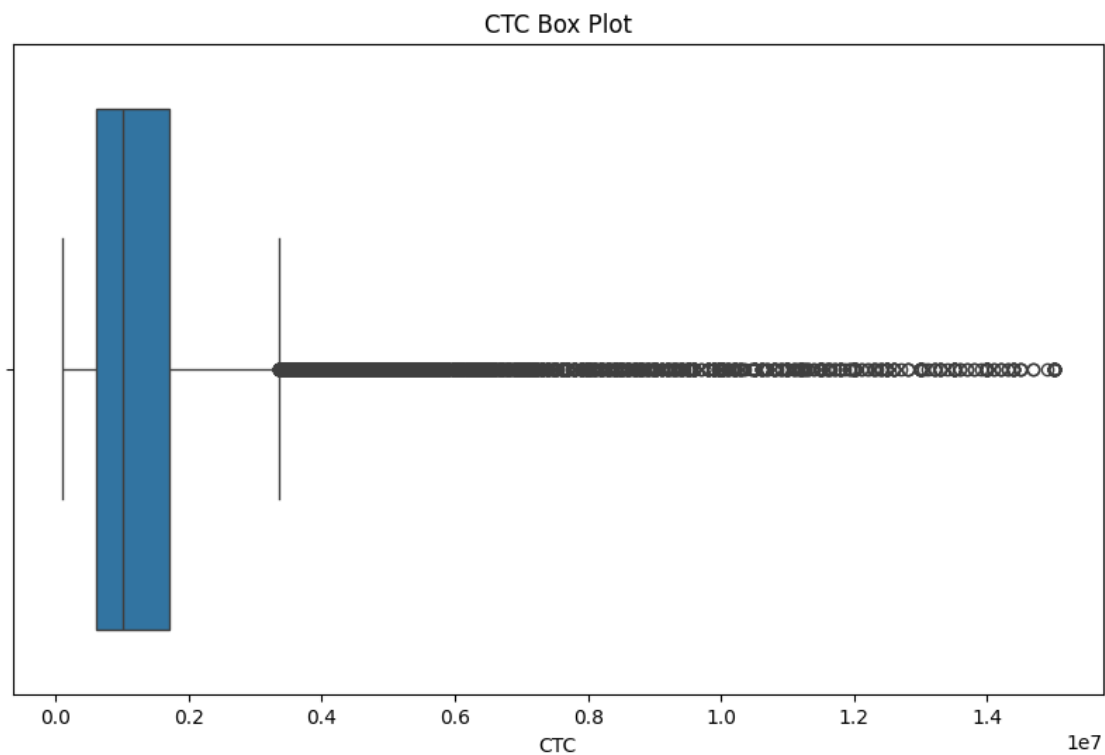
```
[36]: df.groupby('ctc_category')['ctc'].describe()
```

```
[36]:
```

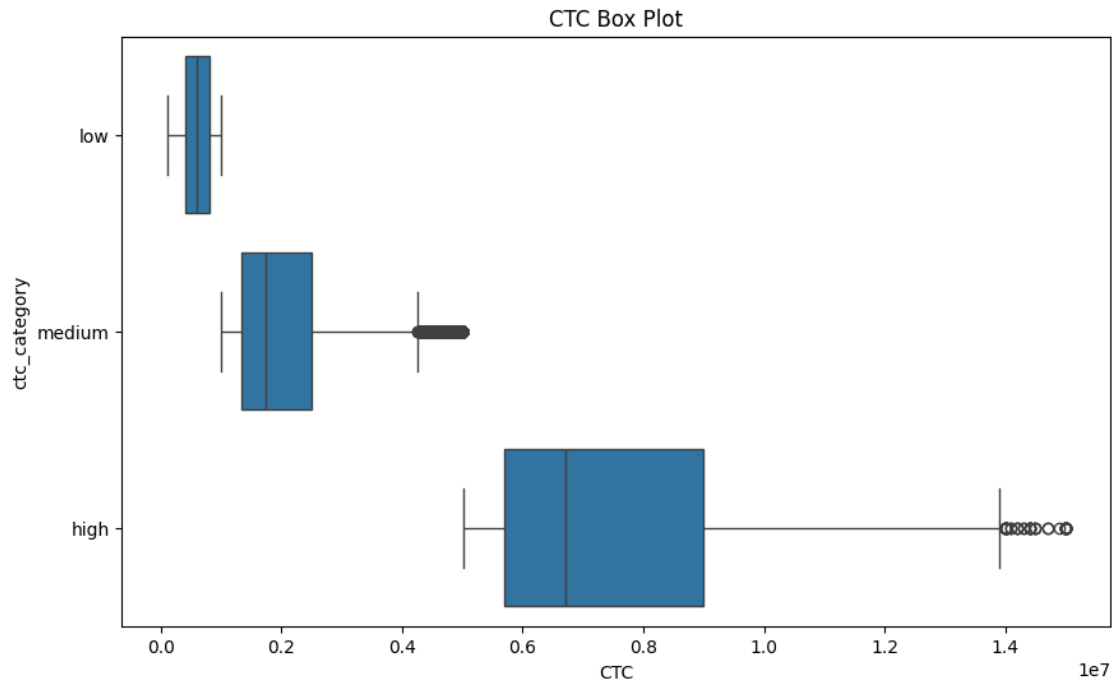
	count	mean	std	min	25% \
ctc_category					
low	91187.0	5.955088e+05	2.401585e+05	100000.0	400000.0
medium	78056.0	2.017507e+06	8.878268e+05	1000008.0	1330000.0
high	3277.0	7.507861e+06	2.281781e+06	5010000.0	5700000.0

	50%	75%	max
ctc_category			
low	600000.0	800000.0	1000000.0
medium	1730000.0	2500000.0	5000000.0
high	6700000.0	9000000.0	15000000.0

```
[37]: plt.figure(figsize=(10, 6))
sns.boxplot(x=df['ctc'])
plt.title('CTC Box Plot')
plt.xlabel('CTC')
plt.show()
```



```
[38]: plt.figure(figsize=(10, 6))
sns.boxplot(x=df['ctc'], y=df['ctc_category'])
plt.title('CTC Box Plot')
plt.xlabel('CTC')
plt.show()
```



```
[39]: print("Organisation Year Range: ", df['orgyear'].min(), df['orgyear'].max())

# Potential outliers - orgyear < 1970 people are not likely to be working after 54 years
df[(df['orgyear'] < 1970) | (df['orgyear'] > 2024)].shape
```

Organisation Year Range: 0 20165

```
[39]: (77, 8)
```

```
[40]: df = df[(df['orgyear'] >= 1970) & (df['orgyear'] <= 2024)]
df['year_of_experience'] = 2024 - df['orgyear']
```

0.1.9 Data Pre-processing:

- **Job Position:** The job title are cleaned manually and mapped backed to the dataset, Missing job positions are filled using KNN Imputation
- **Company Hash:** Dropped the columns where company name not provided

- **CTC:** The dataset has been capped based on the domain knowledge (RS1,00,000 to 1,50,00,000)
- **Organisation Year:** The Organisation year has been capped based on the fact that person might not work after 54 years
- **Year of Experience:** The column derived from the org year

0.1.10 Insights:

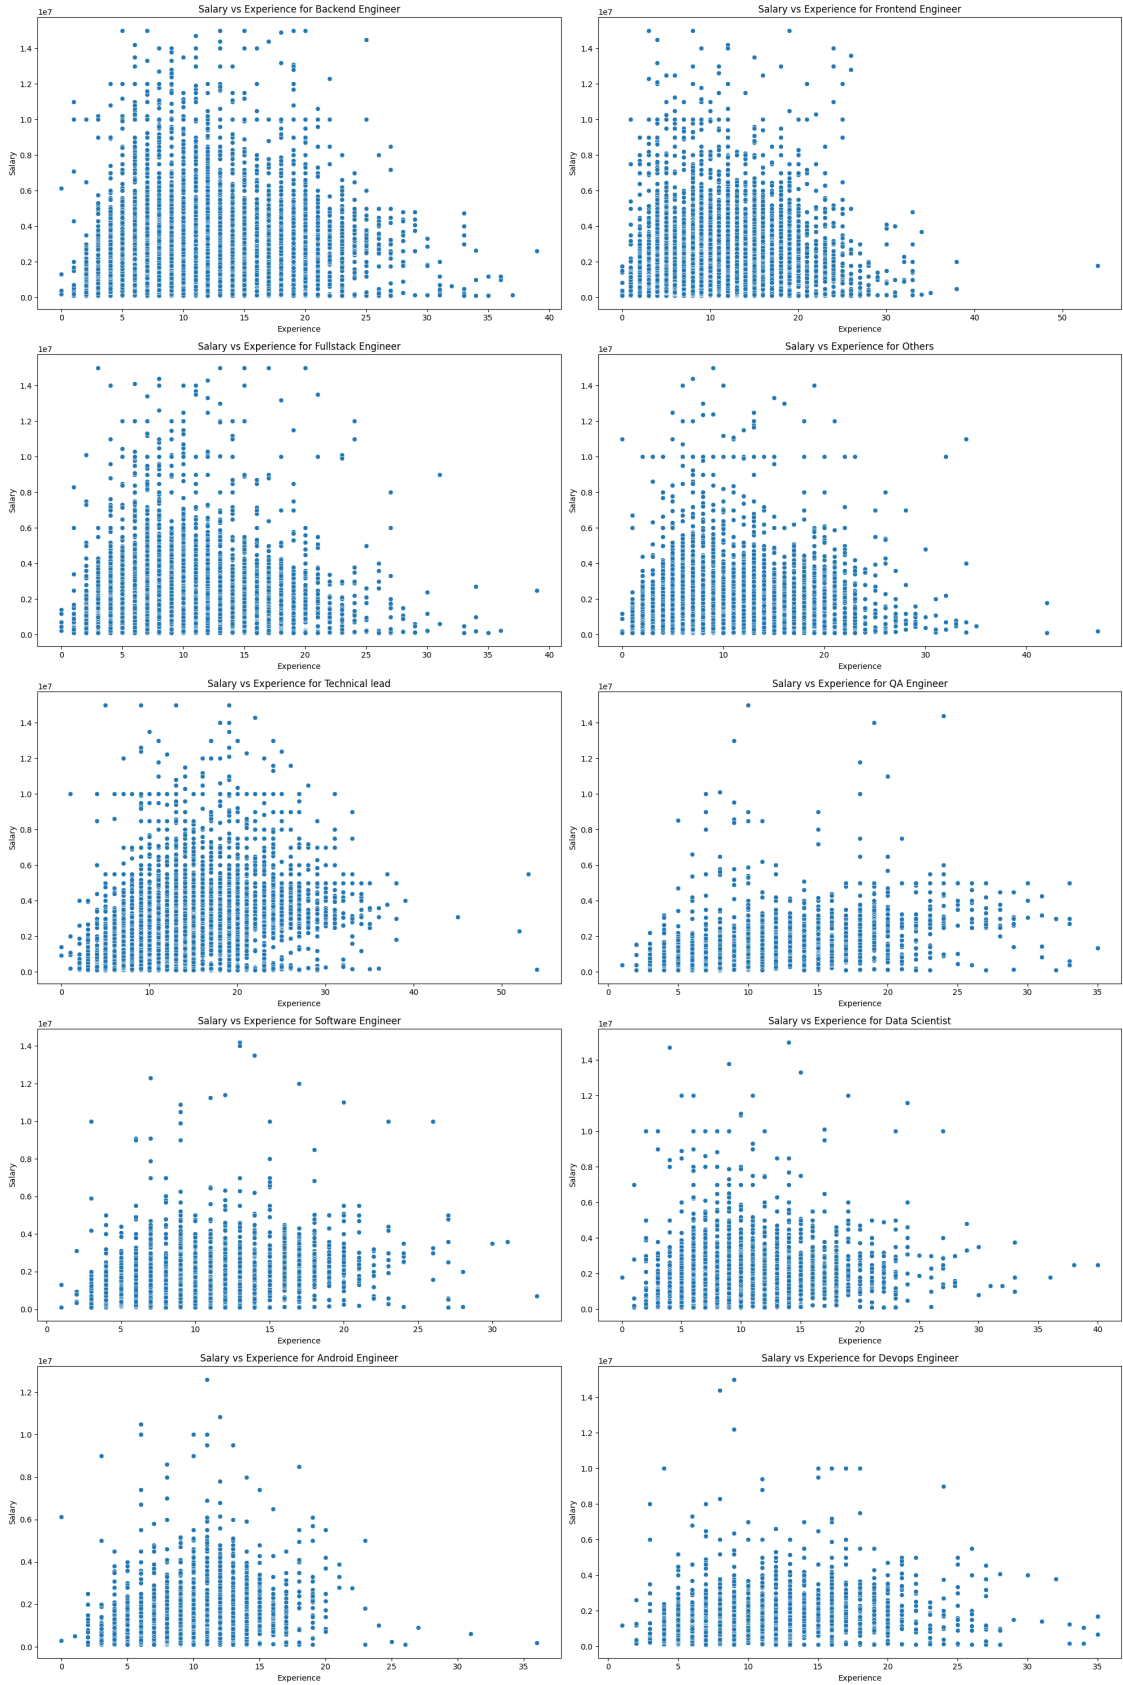
- **Distribution of ctc_category**
 - Low CTC Category: 91,187 learners
 - Medium CTC Category: 78,056 learners
 - High CTC Category: 3,277 learners
- **Skewed Distribution:**
 - The distribution of learners across the ctc_category is highly skewed, with the majority falling into the low and medium categories. Only a small fraction of learners are in the high category.
- **Outliers:**
 - The presence of outliers in the medium and high categories suggests that there are learners with significantly higher CTC values than the rest. These outliers could be due to promotions, exceptional performance, or high-paying job profiles.

```
[41]: # Top 10 job positions, plot the salary vs experience
top_10_job_positions = df['job_position'].value_counts().head(10).index.
↳to_list()

fig, axes = plt.subplots(5, 2, figsize=(20, 30))
axes = axes.flatten()

for i, job in enumerate(top_10_job_positions):
    sns.scatterplot(data=df[df['job_position'] == job], x='year_of_experience', y='ctc', ax=axes[i])
    axes[i].set_title(f'Salary vs Experience for {job}')
    axes[i].set_xlabel('Experience')
    axes[i].set_ylabel('Salary')

plt.tight_layout()
plt.show()
```



0.1.11 Insights:

- **Positive Correlation:**

- There is a general trend where salary increases with years of experience across most job titles. However, this trend is not strictly linear and shows significant variability.

- **Plateau Effect:**

- For many job titles, the salary seems to plateau after a certain number of years of experience. This indicates that beyond a certain point, additional years of experience do not significantly increase the salary.

- **Outliers and High Earners:**

- High-salary outliers are present in most job titles, indicating that exceptional performance, specialized skills, or high-paying employers can lead to significantly higher salaries.

- **Is it always true that with an increase in years of experience, the CTC increases?**

- There is no clear trend, that CTC increase based on increase in years of experience

```
[42]: company_wise_summary_df = df.groupby(['company_hash', 'job_position',  
      ↪ 'year_of_experience'])['ctc'].agg(['min', 'mean', 'median', 'max', 'count']).  
      ↪ reset_index()  
company_wise_summary_df =  
      ↪ company_wise_summary_df[company_wise_summary_df['count'] > 0]  
company_wise_summary_df.sample(5)
```

```
[42]:
```

	company_hash	job_position	\
63643470	wvqo24 otqcxwto uqxcvnt rxbxnta	Backend Engineer	
26039978	ntqvavnv	Research Engineer	
54856399	vngo ojzntr ucn rna	Data Scientist	
46521244	tdutqxnton	Frontend Engineer	
57749595	vuurt	Fullstack Engineer	

	year_of_experience	min	mean	median	max	\
63643470	16	4500000.0	4500000.0	4500000.0	4500000.0	
26039978	6	900000.0	900000.0	900000.0	900000.0	
54856399	17	1800000.0	1800000.0	1800000.0	1800000.0	
46521244	7	2200000.0	2200000.0	2200000.0	2200000.0	
57749595	8	220000.0	220000.0	220000.0	220000.0	

	count
63643470	1
26039978	1
54856399	1

```
46521244      1
57749595      1
```

0.1.12 Feature Creation

```
[43]: def category_mapping(ctc, mean):
        if ctc < mean:
            return 1
        elif ctc > mean:
            return 3
        else:
            return 2
```

```
[44]: df_grouped = df.groupby(['company_hash', 'job_position',
    ↪ 'year_of_experience'])['ctc'].mean().reset_index()
df_grouped.rename(columns={'ctc': 'com_pos_exp_ctc_mean'}, inplace=True)
df_grouped = df_grouped[df_grouped['com_pos_exp_ctc_mean'] > 0]
df_grouped.head()
```

```
[44]:
```


	company_hash	job_position	year_of_experience	\
1180	0	Others	4	
3192	0000	Others	7	
4222	01 ojztsj	Android Engineer	8	
4864	01 ojztsj	Frontend Engineer	13	
6375	05mz exzytvrny uqxcvnt rxbxnta	Backend Engineer	5	

	com_pos_exp_ctc_mean
1180	100000.0
3192	300000.0
4222	270000.0
4864	830000.0
6375	1100000.0

```
[45]: df = pd.merge(df, df_grouped, on=['company_hash', 'job_position',
    ↪ 'year_of_experience'], how='left')
df['designation'] = df.apply(lambda x: category_mapping(x['ctc'],
    ↪ x['com_pos_exp_ctc_mean']), axis=1)
df.drop(columns=['com_pos_exp_ctc_mean'], inplace=True)
```


```
[46]: sns.countplot(x='designation', data=df)
```

```
[46]: <Axes: xlabel='designation', ylabel='count'>
```

```
[47]: sns.boxplot(x='designation', y='ctc', data=df)
```

```
[47]: <Axes: xlabel='designation', ylabel='ctc'>
```



```
[48]: df_grouped = df.groupby(['company_hash', 'job_position'])['ctc'].mean().
      ↪reset_index()
df_grouped.rename(columns={'ctc': 'com_pos_ctc_mean'}, inplace=True)
df_grouped = df_grouped[df_grouped['com_pos_ctc_mean'] > 0]
df_grouped.head()
```

```
[48]:
```

	company_hash	job_position	com_pos_ctc_mean
24	0	Others	100000.0
65	0000	Others	300000.0
86	01 ojztsj	Android Engineer	270000.0
99	01 ojztsj	Frontend Engineer	830000.0
130	05mz exzytvrny uqxcvnt rxbxnta	Backend Engineer	1100000.0

```
[49]: df = pd.merge(df, df_grouped, on=['company_hash', 'job_position'], how='left')
df['Class'] = df.apply(lambda x: category_mapping(x['ctc'],
      ↪x['com_pos_ctc_mean']), axis=1)
df.drop(columns=['com_pos_ctc_mean'], inplace=True)
```

```
[50]: df_grouped = df.groupby(['company_hash'])['ctc'].mean().reset_index()
df_grouped.rename(columns={'ctc': 'com_ctc_mean'}, inplace=True)
```

```
df_grouped = df_grouped[df_grouped['com_ctc_mean'] > 0]
df_grouped.head()
```

```
[50]:
```

	company_hash	com_ctc_mean
0	0	100000.0
1	0000	300000.0
2	01 ojztqsj	550000.0
3	05mz exzytvrny uqxcvnt rxbxnta	1100000.0
4	1	175000.0

```
[51]: df = pd.merge(df, df_grouped, on=['company_hash'], how='left')
df['Tier'] = df.apply(lambda x: category_mapping(x['ctc'], x['com_ctc_mean']),
axis=1)
df.drop(columns=['com_ctc_mean'], inplace=True)
```

```
[52]: df.sample(5)
```

```
[52]:
```

	company_hash \
87394	ztdngqj uqxcvnt rxbxnta
74538	tdwtrrtd ntwyzgrgsxto uqxcvnt rxbxnta
27458	ti ntwyzgrgsxw
42287	ojzguojo xzw
35043	zvz


	email_hash	orgyear	ctc \
87394	81f9e9d24d143e18a6b4e4ea64792038392ca1d5f32f5a...	2019	550000
74538	6ed3aa0728ec2819fd74bb9a56e94121e9affa4ed238ce...	2015	672000
27458	286e211d9ac97b8f644751c955614bdac0056f4bc3cea0...	2017	1500000
42287	3ea8a32a8fa5f84afbb931771154609428bb29065d00d3...	2020	380000
35043	33e1fc8bd06fc0f20c643df27a2667bcad34d2d6ba66fb...	2015	2200000

	job_position	ctc_updated_year	CTC_zscore	ctc_category \
87394	Frontend Engineer	2021	-0.149401	low
74538	Fullstack Engineer	2021	-0.139541	low
27458	Frontend Engineer	2021	-0.072625	medium
42287	Backend Engineer	2019	-0.163140	low
35043	Frontend Engineer	2021	-0.016053	medium

	year_of_experience	designation	Class	Tier
87394	5	2	2	2
74538	9	2	2	2
27458	7	3	3	3
42287	4	1	1	1
35043	9	3	3	3

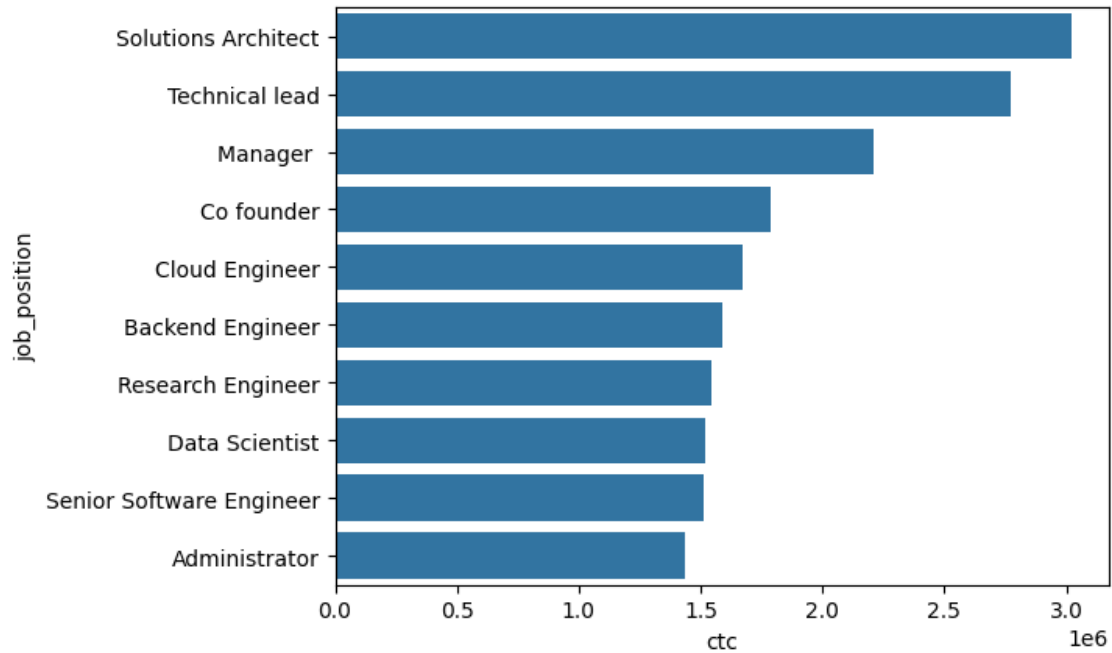
```
[53]: sns.barplot(x='ctc_category', y='year_of_experience', data=df)
```

[53]: <Axes: xlabel='ctc_category', ylabel='year_of_experience'>



```
[54]: # Top 10 Job positions with highest mean CTC
top_10_job_positions = df.groupby('job_position')['ctc'].mean().
    ↪sort_values(ascending=False).head(10).reset_index()
sns.barplot(x='ctc', y='job_position', data=top_10_job_positions)
```

[54]: <Axes: xlabel='ctc', ylabel='job_position'>



0.1.13 Insights:

Top 10 Job positions with highest mean CTC * Solutions Architect * Technical lead * Manager * Co founder * Cloud Engineer * Backend Engineer * Research Engineer * Data Scientist * Senior Software Engineer * Administrator

```
[55]: df[df['Tier']==1].sort_values(by='ctc', ascending=False).head(10)
```

```
[55]:
```

	company_hash \		email_hash	orgyear	ctc \
70681	aggqavoy		68f1fea4dbfb7ae2209664b93d5f57fb86912dbe516b37...	2018	13500000
62401	mqxwpontr tzntquqxoto uqxcvnt rxbxnta		5cb0417e963b2bb218dc28cbe0c9003c39c4f2db94bb53...	2016	10100000
44445	bsb qtogqno xzntqzvnxyzvr		420388fd953332be671e1b0761f9af06d323382d075ecf...	2017	7000000
31001	zvnxyzvr vhoneqrvxv mvzp		2ddbc233754a1bf09fa7e92d61a5fb8fd46f3fe7908318...	2000	7000000
42591	xzagqot unt rna		3f0f3d507158974626454f18f97876cd1f3ffe816a9757...	2008	6500000
141931	wgzerhtzn				
118650	uqgutqnjshqh				
17465	wvqttb				
103987	uvd vx				
103730	xatvrg				

141931	d2baa62a1b611ad438f0475c2e3d88f2c6f9f3df6e56bf...	2015	6000000
118650	b01a6b018bb4ce949e745045fae952bd9c89566480e066...	2015	5820000
17465	19839c1100097c37582a06a4dbb8ec28194883f0c6b70f...	2017	5600000
103987	9a6853551e4a153f057d06ae57bfba4a7027cdf43cf273...	2012	5600000
103730	9a060e939ffc6ec3f5a9351ffb303ad6d064761ee66c94...	2000	5440000

	job_position	ctc_updated_year	CTC_zscore	ctc_category \
70681	Backend Engineer	2020	0.897179	high
62401	QA Engineer	2019	0.622401	high
44445	Fullstack Engineer	2018	0.371868	high
31001	Technical lead	2020	0.371868	high
42591	Backend Engineer	2020	0.331460	high
141931	Backend Engineer	2020	0.291051	high
118650	Fullstack Engineer	2020	0.276504	high
17465	Backend Engineer	2020	0.258725	high
103987	Backend Engineer	2019	0.258725	high
103730	Technical lead	2020	0.245794	high

	year_of_experience	designation	Class	Tier
70681	6	2	2	1
62401	8	1	1	1
44445	7	2	2	1
31001	24	2	2	1
42591	16	1	1	1
141931	9	2	2	1
118650	9	2	1	1
17465	7	1	1	1
103987	12	2	2	1
103730	24	2	2	1

```
[56]: # Top 10 employees of data science in each company earning more than their
      ↪peers - Class 1
top_10_employees = df[(df['job_position']=='Data Scientist') & (df['Class']==1)]
top_10_employees = top_10_employees.groupby('company_hash').apply(lambda x: x.
      ↪nlargest(10, 'ctc')).reset_index(drop=True)
top_10_employees.sort_values(by='company_hash')
```

```
[56]: company_hash \
0      1bs
1      247vx
2      247vx
3      3p ntwyzgrgsxto
4      3p ntwyzgrgsxto
...      ...
1491     zxxn ntwyzgrgsxto
1493  zxxn ntwyzgrgsxto rxbxnta
1495     zxztrtvuo
```

```
1494          zxztrtvuo
1496          zxztrtvuo
```

```

                                email_hash  orgyear      ctc  \
0      eb213c0552effd7fb139395c7838edb8d59773a1cb57a0...  1994  800000
1      c35054c043f6a02da3e6f142fbc095f8145eb521137ff...  2014  2150000
2      5f4b52a1c2539fe2e4b29a8470bc57dbace331b819a0af...  2002  1440000
3      583a9600d8e74d5f3f6627da6b4ff3c466bf5cfee5ae9f...  2018  1200000
4      f3b1e96456ad8a7d9cceb3901d763252ea2f56eb2ee7f4...  2013  1200000
...
1491  1aa2717970a46b5d12b90932799227774dd418c842fa18...  2012  2200000
1493  5ab93fd511bceaa6da5f855d160de306a04df9951f5978...  2012   800000
1495  3027ca561b65f99da2f65bf3d85c6bb5d5687c67e69e89...  2018  1370000
1494  10d566c5fca40ffe1d133b79594d071880711ef480da9f...  2017  1400000
1496  f678c67bee8cad9370f6aaf4f4cc22ffd417fd753663c6...  2019  1250000
```

```

      job_position  ctc_updated_year  CTC_zscore  ctc_category  \
0      Data Scientist          2019    -0.129197          low
1      Data Scientist          2018    -0.020094        medium
2      Data Scientist          2019    -0.077474        medium
3      Data Scientist          2019    -0.096870        medium
4      Data Scientist          2019    -0.096870        medium
...
1491  Data Scientist          2019    -0.016053        medium
1493  Data Scientist          2021    -0.129197          low
1495  Data Scientist          2019    -0.083131        medium
1494  Data Scientist          2019    -0.080707        medium
1496  Data Scientist          2021    -0.092829        medium
```

```

      year_of_experience  designation  Class  Tier
0              30              2      1      1
1              10              2      1      3
2              22              2      1      1
3               6              1      1      3
4              11              2      1      3
...
1491              12              2      1      3
1493              12              2      1      1
1495               6              2      1      3
1494               7              2      1      3
1496               5              2      1      3
```

[1497 rows x 12 columns]

[57]:

```

# Bottom 10 employees of data science in each company earning less than their
↳peers - Class 3
bottom_10_employees = df[(df['job_position']=='Data Scientist') &
↳(df['Class']==3)]

# Initialize an empty DataFrame to store the results
result_df = pd.DataFrame()

# Process each company separately to avoid memory issues
for company in bottom_10_employees['company_hash'].unique():
    company_df = bottom_10_employees[bottom_10_employees['company_hash'] ==
↳company]
    bottom_10 = company_df.groupby('email_hash')['ctc'].mean().
↳nsmallest(10).reset_index()
    bottom_10['company_hash'] = company
    result_df = pd.concat([result_df, bottom_10], ignore_index=True)

result_df = result_df[result_df['ctc'].notnull()]
result_df

```

```

[57]:

```

	email_hash	ctc \
0	63ea6962348c633656fb3e335c41356c524ed6559176a7...	1610000.0
1	001f9dbae7fccf1ef52c2187359d65c17e2897b8d211d2...	2000000.0
2	a59dd7dfe78bdb896a0aa81b0fa938cd4459a99f14fa03...	2000000.0
10	0038257f6ef2580dfbc8566cc80d519819c97c2dcbddde...	4200000.0
20	70e3f9611e8afdd1184c174d0948575c6fc3dba0b1aa21...	1050000.0
...
6340	fb5db1053d6c3d953112bd6316b9d0ebe2954dc4df3c78...	3500000.0
6350	fc0a94acd05743069d35b4dab69efd913fce055f374c8f...	2500000.0
6360	fc48d558c7d9abc5b8ca58c524173255b26782c15c28b9...	2000000.0
6370	fd97ee5909e0db824aaf15b5f293987a1c21272ecc491e...	5600000.0
6380	ffb56190424c95d9b895e9e3712cdc94c4382011357d18...	1100000.0
	company_hash	
0	xzwtag xzw	
1	xzwtag xzw	
2	xzwtag xzw	
10	rn ntwyzgrgsj otqcxwto rxbxnta	
20	tj	
...	...	
6340	ovrtoegqwt	
6350	ktzaqxct	
6360	ihxpq	
6370	jghurho xzw	
6380	vrsgzgd ntwyzgrgsxto	

[1261 rows x 3 columns]


```
[58]: # Bottom 10 employees (earning less than most of the employees in the company)
      ↪- Tier 3
bottom_10_employees = df[df['Tier']==3]

# Initialize an empty DataFrame to store the results
result_df = pd.DataFrame()

# Process each company separately to avoid memory issues
for company in bottom_10_employees['company_hash'].unique():
    company_df = bottom_10_employees[bottom_10_employees['company_hash'] ==
    ↪company]
    bottom_10 = company_df.groupby('email_hash')['ctc'].mean().
    ↪nsmallest(10).reset_index()
    bottom_10['company_hash'] = company
    result_df = pd.concat([result_df, bottom_10], ignore_index=True)

result_df = result_df[result_df['ctc'].notnull()]
result_df.sort_values(by='ctc', ascending=True).head(10)
```

```
[58]:
```

	email_hash	ctc \
82840	df98a0f24fb45172e2768b53515d980f2ff5b40265d8b8...	105000.0
65050	979c186a6cd175b2dfbfea8e50d3bb823bb4fb9d2fd5e5...	110000.0
33330	369ccf8ed6bf03c964d2822f76119ba94776b0c9ad7485...	115000.0
15510	110f894dcc2dea3cbdcdbd8f6e5857407184c9062920a30...	115000.0
75360	c0e59bea31301b6f14bcb8bf8ff87485ffdeac52a5d5f3...	120000.0
900	00972ad5cc152a8763c501413f1213b4c0c9cb78ed735b...	120000.0
83090	e12764654b74f407aa9ddf4819030c1d9cc88e9c3f359f...	120000.0
33331	9c8846c710a0266e74dd7d29e0d3d61f533a0da533e1ac...	120000.0
75600	c1b86ff295f82151e1649a338eaae941766375a74fca96...	120000.0
78290	cc51b85d8792b7755d107504dfa1ffa6123ade28c65bd7...	120000.0

	company_hash
82840	eqjo trtnwqgzxwo
65050	ctqxkgz fxqtrtoo
33330	oyhnntqerj xzw
15510	wxqwrto rxet
75360	uvooqh vtqgouvwt xzw
900	ovzaxv zvnxgzvr rvmgqvngqxto
83090	uwo srgmvr uqxcvnt rxbxnta
33331	oyhnntqerj xzw
75600	xytvqnqvaxg
78290	tuxw svbto

```
[59]: # Top 10 employees in each company - X department - having 5/6/7 years of
      ↪experience earning more than their peers - Tier X

# Initialize an empty DataFrame to store the results
```

```

result_df = pd.DataFrame()

# Filter the relevant data
filtered_df = df[
    (df['job_position']=='Data Scientist') & (df['year_of_experience'].
    ↪isin([5,6,7])) & (df['Tier']==1)
]

# Process each company separately to avoid memory issues
for company in filtered_df['company_hash'].unique():
    company_df = filtered_df[filtered_df['company_hash'] == company]
    top_10 = company_df.groupby('email_hash')['ctc'].mean().nlargest(10).
    ↪reset_index()
    top_10['company_hash'] = company
    result_df = pd.concat([result_df, top_10], ignore_index=True)

result_df = result_df[result_df['ctc'].notnull()]
result_df.sort_values(by='ctc', ascending=True).head(10)

```

```

[59]:

```

	email_hash	ctc \
2980	6d50e60d562940b28dc7cd0296ea52164d272d7f5696f8...	100000.0
3321	7ff23b1d29cb8a2546289c567848c6f2ebbcd335af9f5c...	100000.0
1317	2cadff2fe2f39e9e4e684b95c0db58c9344fceb815c6c1...	100000.0
1820	3f1cd17776c41c8395107a340b76af3a56fe307987e271...	103000.0
400	0a65cc19cc9402203e3282318b9a4a127eb0b6691120cc...	110000.0
1430	30b64da7d7e52c133a0e1df85c912794081242e0a792fa...	110000.0
5360	fa24a6efde82936d129f71557a1723fcefcf0db9eee432...	110000.0
3320	a9560804241d30092205375646cd9c0df1c4a49204e678...	112000.0
4800	db2c70fea469a7f1456457812fe94a01c337eb6ce75bd5...	115000.0
1670	38b81aab25e388fb3282b99e6872365cf15b80a08bd787...	120000.0

	company_hash
2980	ktgnvu
3321	ytrrgoxcx ogenfvqt rvmo
1317	zgzt
1820	hztburgjta
400	ihvznuyx
1430	vhaxmrt xzw
5360	xzatrxtzn
3320	ytrrgoxcx ogenfvqt rvmo
4800	mvjqtq
1670	xzwgqnv

```

[60]: # Initialize an empty DataFrame to store the results
result_df = pd.DataFrame()

# Filter the relevant data

```

```

filtered_df = df[
    (df['job_position']=='Data Scientist') & (df['year_of_experience'].
    ↪isin([5,6,7])) & (df['Tier']==3)
]

# Process each company separately to avoid memory issues
for company in filtered_df['company_hash'].unique():
    company_df = filtered_df[filtered_df['company_hash'] == company]
    top_10 = company_df.groupby('email_hash')['ctc'].mean().nlargest(10).
    ↪reset_index()
    top_10['company_hash'] = company
    result_df = pd.concat([result_df, top_10], ignore_index=True)

result_df = result_df[result_df['ctc'].notnull()]
result_df

```

```

[60]:

```

		email_hash	ctc \
0		009ded427ebcb5c2fb1970017a683693a7abef0fa96f5e...	3900000.0
1		ea7eb74cbbabcb83aa2f66891c6deb847a603c2a40535f...	2800000.0
2		10dc906ae6fedaac7e16b0dd71356be39e2136911990a9...	2380000.0
3		3fdfb9b1014a0d22fd5fa9d5c7a686a31a16a933eda144...	2300000.0
4		aedf693542ac76087cb458cfc66154382568d02f0d6acb...	2200000.0
...	
3890		fba008a8a93022172e3adfa293d9ce0f4d70267afa373d...	500000.0
3900		fbac8d7a3768fdb72d929a641d9ec1c307d35487bd0aad...	1500000.0
3910		fbb5702400e9542c9ddcbb561189dafec71131f8458856...	1050000.0
3920		fc48d558c7d9abc5b8ca58c524173255b26782c15c28b9...	2000000.0
3930		ffc974693a2bfd0326c707d8460d6783861a9497e538e2...	1700000.0

```


```

		company_hash
0		eqvwnvr vzvrjnxwo
1		eqvwnvr vzvrjnxwo
2		eqvwnvr vzvrjnxwo
3		eqvwnvr vzvrjnxwo
4		eqvwnvr vzvrjnxwo
...		...
3890		vuurxta vx
3900	bvyxzaqv	wgbcxcv ntwyzgrgsxto rna
3910		bvqgugon
3920		ihxpq
3930		atrgxnnt xzaxv

```

[566 rows x 3 columns]

```

```

[61]: # Initialize an empty DataFrame to store the results
result_df = pd.DataFrame()

```

```

# Filter the relevant data
filtered_df = df[
    (df['job_position']=='Data Scientist') & (df['year_of_experience'].
    ↪isin([5,6,7])) & (df['Tier']==2)
]

# Process each company separately to avoid memory issues
for company in filtered_df['company_hash'].unique():
    company_df = filtered_df[filtered_df['company_hash'] == company]
    top_10 = company_df.groupby('email_hash')['ctc'].mean().nlargest(10).
    ↪reset_index()
    top_10['company_hash'] = company
    result_df = pd.concat([result_df, top_10], ignore_index=True)

result_df = result_df[result_df['ctc'].notnull()]
result_df

```

```

[61]:

```

	email_hash	ctc \
0	013a80ea3b5799eb374ec5dd94d3c84579ba0b44293258...	400000.0
10	01e291d69add1b6340ff19d8f4cb396ffb2a7a58c5b82b...	900000.0
20	01e29469795e1395f058e33226e15b1a32d335e92ac5bf...	1800000.0
30	0272dcd1909235f14140851a86a1bfb22098d4c71bd0f7...	1500000.0
40	02d70fdfcde30937fa1953b7e17b4f20b6bb57de684ec5...	700000.0
...
3600	fa57641000b5bb092b93eb2d5b615275b550045aece5b9...	750000.0
3610	fb62ad17e33d701e2369045284190ff75e7e795eba1dc5...	620000.0
3620	fc0472f7b02f8569d8a368972f2f6f99e509a961e9078e...	1980000.0
3630	ff7bace89454b0965e40533aa5b7f904dba9c496b1a1f2...	1019999.0
3640	ffa4d3df725be3628607627e575feac064a1e98506dc19...	900000.0

	company_hash
0	mxsnvuu vzvrjnxwo
10	avnv owxtzwt nqvzxtt vn tafxogq
20	tuxex ntwyzgrgsxto ucn rna
30	vqjv
40	vqjzs
...	...
3600	bvszv xzegntwy v ihtoo wgbuvzj
3610	vae avnv owxtzwt uqxcvnt rxbxnta
3620	l u bgqsvz otqcxwto xzaxv ucn rna
3630	xuqtaxwnwgszxnxt ntwyzgrgsxto ucn rna
3640	qtmxn

[366 rows x 3 columns]

```

[62]:

```

```
# Top 10 companies (based on their CTC)
top_10_companies = df.groupby('company_hash')['ctc'].mean().
↳sort_values(ascending=False).head(10).reset_index()
top_10_companies
```

```
[62]:
```

	company_hash	ctc
0	mqvojo	15000000.0
1	qtuogr	15000000.0
2	ihxu	15000000.0
3	cyhm	15000000.0
4	qtvrotre	15000000.0
5	hzxctqoxnj ge ntdvo vn avrrvo	15000000.0
6	vnox xzw	14400000.0
7	ogenfvqt tzsxzttqxzs	14400000.0
8	gqvmvot ogrhnxgzo	14400000.0
9	xnc	14000000.0

```
[63]: # Top 2 positions in every company (based on their CTC)
top_2_positions = df.groupby(['company_hash', 'job_position'])['ctc'].mean().
↳groupby('company_hash', group_keys=False).nlargest(2).reset_index()
top_2_positions.groupby('company_hash')['job_position'].apply(list).
↳reset_index()
```

```
[63]:
```

	company_hash	job_position
0	0	[Others, Accounting]
1	0000	[Others, Accounting]
2	01 ojztsj	[Frontend Engineer, Android Engineer]
3	05mz exzytvrny uqxcvnt rxbxnta	[Backend Engineer, Accounting]
4	1	[Others, Frontend Engineer]
...
37294	zyvzwt wgzohrnxs tzsxzttqo	[Frontend Engineer, Accounting]
37295	zz	[Others, Frontend Engineer]
37296	zzb ztdnstz vacxogqj ucn rna	[Fullstack Engineer, Accounting]
37297	zzgato	[Frontend Engineer, Accounting]
37298	zzzbzb	[Others, Accounting]

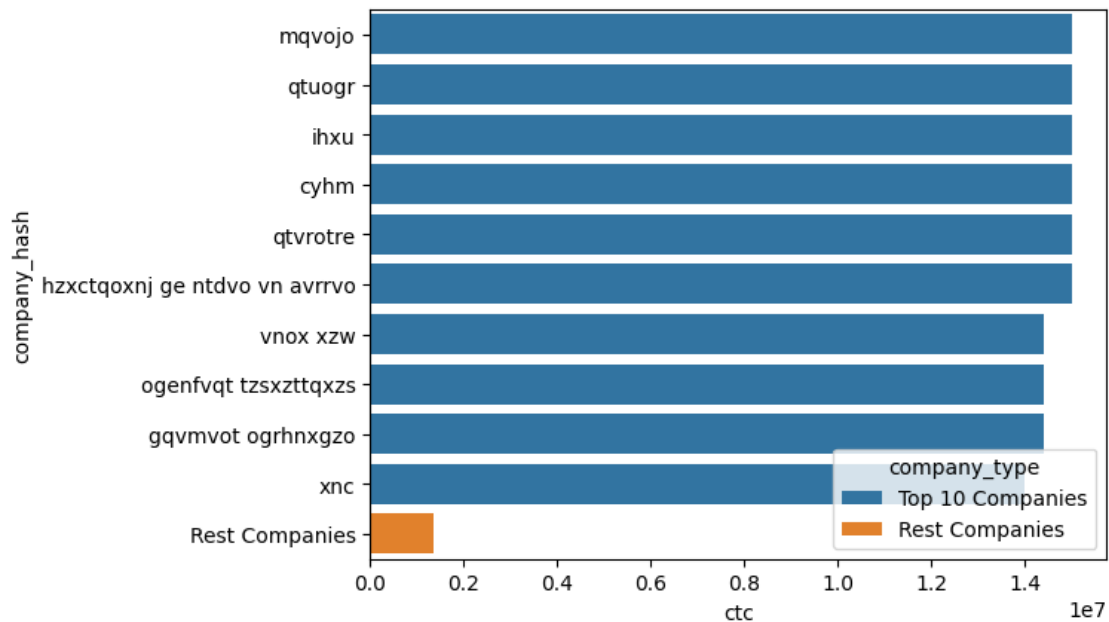
[37299 rows x 2 columns]

```
[64]: # compare the mean CTC of the top 10 companies with the mean CTC of the rest of
↳the companies
top_10_companies = df.groupby('company_hash')['ctc'].mean().
↳sort_values(ascending=False).head(10).reset_index()
top_10_companies['company_type'] = 'Top 10 Companies'

rest_companies = df[~df['company_hash'].isin(top_10_companies['company_hash'])]
rest_companies_mean_ctc = rest_companies['ctc'].mean()
```

```
rest_companies = pd.DataFrame({'company_hash': ['Rest Companies'], 'ctc':  
    ↳[rest_companies_mean_ctc], 'company_type': ['Rest Companies']})  
  
companies_df = pd.concat([top_10_companies, rest_companies], ignore_index=True)  
sns.barplot(x='ctc', y='company_hash', data=companies_df, hue='company_type')
```

[64]: <Axes: xlabel='ctc', ylabel='company_hash'>



[65]: df.sample(5)

```
[65]:
```

	company_hash \							
150583	kxusg							
112825	rtznqv							
62856	exwg							
25746	uqvbvnx ntwyzgrgsxto uqxcvnt rxbxnta							
133659	whqgej							

	email_hash	orgyear	ctc \
150583	dfa84e64857e65fb415417939f6393f520462cae373a4f...	2014	1700000
112825	a787c58d015cfb352a1d4244d52da45738169326b999f2...	2018	660000
62856	5d6d61c7f0a99ff6fc574b8ed0263eb3132d060cfc575e...	2010	2000000
25746	25dfdf2bbb2b5f45fb5f82ed2d7620f7b25806431df906...	2014	600000
133659	c682f7df08f21debe505a5eaca754e05b79b5a6a17e060...	2010	700000

	job_position	ctc_updated_year	CTC_zscore	ctc_category \
150583	Manager	2019	-0.056461	medium

112825	Backend Engineer	2021	-0.140511	low
62856	Backend Engineer	2019	-0.032216	medium
25746	QA Engineer	2019	-0.145360	low
133659	Frontend Engineer	2019	-0.137278	low

	year_of_experience	designation	Class	Tier
150583	10	2	1	3
112825	6	2	2	3
62856	14	3	3	3
25746	10	2	2	1
133659	14	2	1	1

0.1.14 Data preparation for Modelling

```
[74]: scaled_df = df.copy()
```

```
[75]: # Standardize Numerical Data
numerical_cols = ['ctc', 'year_of_experience', 'orgyear', 'ctc_updated_year']
scaler = StandardScaler()
scaled_df[numerical_cols] = scaler.fit_transform(scaled_df[numerical_cols])

# Frequency Encoding for company_hash
company_hash_frequency = scaled_df['company_hash'].value_counts().to_dict()
scaled_df['company_hash_encoded'] = scaled_df['company_hash'].
    ↪map(company_hash_frequency)

# Encode Other Categorical Data
categorical_cols = ['job_position', 'ctc_category', 'Class', 'Tier']
encoder = OneHotEncoder(sparse_output=False)
encoded_categorical_data = encoder.fit_transform(scaled_df[categorical_cols])
encoded_categorical_df = pd.DataFrame(encoded_categorical_data, columns=encoder.
    ↪get_feature_names_out(categorical_cols))

# Concatenate the encoded categorical data with the original DataFrame
scaled_df = pd.concat([scaled_df, encoded_categorical_df], axis=1)

# Drop the original categorical columns
scaled_df.drop(columns=categorical_cols, inplace=True)
scaled_df.drop(columns=['company_hash', 'CTC_zscore', 'email_hash'], inplace=True)
```

```
[76]: scaled_df.sample(5)
```

```
[76]:      orgyear      ctc  ctc_updated_year  year_of_experience  designation \
128407  0.465676  0.581414          1.115702          -0.465676           3
132177  0.230061 -0.665794          1.115702          -0.230061           2
```

142402	-1.890475	0.252396	-0.388753	1.890475	2
144312	-0.712400	-0.589278	1.115702	0.712400	1
108296	-0.948015	-0.359730	-1.893209	0.948015	3

	company_hash_encoded	job_position_Accounting	\
128407	1459	0.0	
132177	2	0.0	
142402	52	0.0	
144312	41	0.0	
108296	288	0.0	

	job_position_Administrator	job_position_Advisory	\
128407	0.0	0.0	
132177	0.0	0.0	
142402	0.0	0.0	
144312	0.0	0.0	
108296	0.0	0.0	

	job_position_Analyst	...	job_position_iOS Engineer	\
128407	0.0	...	0.0	
132177	0.0	...	0.0	
142402	0.0	...	0.0	
144312	0.0	...	0.0	
108296	0.0	...	0.0	

	ctc_category_high	ctc_category_low	ctc_category_medium	Class_1	\
128407	0.0	0.0	1.0	0.0	
132177	0.0	1.0	0.0	0.0	
142402	0.0	0.0	1.0	0.0	
144312	0.0	1.0	0.0	1.0	
108296	0.0	1.0	0.0	0.0	

	Class_2	Class_3	Tier_1	Tier_2	Tier_3
128407	0.0	1.0	0.0	0.0	1.0
132177	1.0	0.0	0.0	0.0	1.0
142402	0.0	1.0	0.0	0.0	1.0
144312	0.0	0.0	1.0	0.0	0.0
108296	0.0	1.0	1.0	0.0	0.0

[5 rows x 56 columns]


Model Building

KMeans Clustering


```
[77]: X = scaled_df.values


# Elbow Method
wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, random_state=42)
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)

plt.figure(figsize=(10, 6))
plt.plot(range(1, 11), wcss, marker='o')
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.show()
```



```
[78]: # Silhouette Analysis
silhouette_scores = []
for i in range(2, 8):
    kmeans = KMeans(n_clusters=i, random_state=42)
    kmeans.fit(X)
    score = silhouette_score(X, kmeans.labels_)
    silhouette_scores.append(score)
```

```
plt.figure(figsize=(10, 6))
plt.plot(range(2, 8), silhouette_scores, marker='o')
plt.title('Silhouette Analysis')
plt.xlabel('Number of clusters')
plt.ylabel('Silhouette Score')
plt.show()
```



0.1.15 Elbow Method and Silhouette Analysis

- **Elbow Method:** Suggests a potential optimal number of clusters at 3, as there is a noticeable drop in WCSS.
- **Silhouette Analysis:** Indicates that 4 clusters have the highest silhouette score, suggesting the best-defined clusters.


Considering both methods, the optimal number of clusters appears to be 4 clusters. While the Elbow Method suggests 3 clusters, the Silhouette Analysis strongly indicates that 4 clusters provide the best-defined clusters. Therefore, 4 clusters are likely the most appropriate choice for this dataset.

```
[118]: kmeans = KMeans(n_clusters = 4, random_state = 3)
kmeans.fit(X)

scaled_df['kM_cluster'] = kmeans.labels_
scaled_df['kM_cluster'].value_counts()
```

```
[118]: kM_cluster
      1    138028
      3     20824
      2      7189
      0      6402
      Name: count, dtype: int64
```

```
[119]: # Distribution of each cluster
plt.figure(figsize=(10, 6))
scaled_df['kM_cluster'].value_counts().plot(kind='bar')
plt.title('Cluster Distribution')
plt.xlabel('Cluster')
plt.ylabel('Count')
plt.show()
```



```
[120]: import matplotlib.pyplot as plt
from sklearn.decomposition import PCA


# Perform PCA to reduce dimensionality to 2D for visualization
pca = PCA(n_components=2)
pca_result = pca.fit_transform(scaled_df.drop(columns=['kM_cluster']))

# Create a DataFrame with PCA results and cluster labels
pca_df = pd.DataFrame(pca_result, columns=['PCA1', 'PCA2'])
pca_df['Cluster'] = scaled_df['kM_cluster']
```

```

# Plot the PCA results
plt.figure(figsize=(10, 6))
plt.scatter(pca_df['PCA1'], pca_df['PCA2'], c=pca_df['Cluster'],
            cmap='viridis', alpha=0.5)
plt.title('PCA Visualization of Clusters')
plt.xlabel('PCA1')
plt.ylabel('PCA2')
plt.colorbar(label='Cluster')
plt.show()

```



```

[127]: df['kM_cluster'] = kmeans.labels_
print('-'*5, 'Cluster Analysis', '-'*5)
print(df.groupby('kM_cluster')['year_of_experience'].mean())
print('-'*10, 'Cluster Mean', '-'*10)
print(df.groupby('kM_cluster')['ctc'].mean())

```

```

----- Cluster Analysis -----
kM_cluster
0    6.785848
1    9.271662
2    7.678398
3    8.141087
Name: year_of_experience, dtype: float64

```

----- Cluster Mean -----

kM_cluster

0 6.138651e+05

1 1.418803e+06

2 1.366223e+06

3 1.281431e+06

Name: ctc, dtype: float64

```
[128]: print('-'*10, 'Cluster Median', '-'*10)
print(df.groupby('kM_cluster')['ctc'].median())
print('-'*10, 'Cluster Minimum', '-'*10)
print(df.groupby('kM_cluster')['ctc'].min())
print('-'*10, 'Cluster Maximum', '-'*10)
print(df.groupby('kM_cluster')['ctc'].max())
```

----- Cluster Median -----

kM_cluster

0 450000.0

1 1040000.0

2 700000.0

3 819999.0

Name: ctc, dtype: float64

----- Cluster Minimum -----

kM_cluster

0 100000

1 100000

2 100000

3 100000

Name: ctc, dtype: int32

----- Cluster Maximum -----

kM_cluster

0 11200000

1 15000000

2 14700000

3 15000000

Name: ctc, dtype: int32

0.1.16 Interpretation

Cluster 0: - Median CTC: This cluster has the lowest median CTC at 450,000, indicating that the majority of individuals in this cluster have relatively lower salaries compared to other clusters.

- **Range:** The CTC ranges from 100,000 to 11,200,000, showing a significant spread. This suggests that while the median is low, there are some high earners within this cluster.

Cluster 1: - Median CTC: This cluster has the highest median CTC at 1,040,000, indicating that the majority of individuals in this cluster have higher salaries compared to other clusters.

- **Range:** The CTC ranges from 100,000 to 15,000,000, indicating a wide spread. This cluster includes both high earners and individuals with lower salaries.

Cluster 2: - Median CTC: This cluster has a median CTC of 700,000, which is moderate compared to other clusters. - **Range:** The CTC ranges from 100,000 to 14,700,000, showing a wide spread. This cluster includes a mix of individuals with varying salary levels.

Cluster 3: - Median CTC: This cluster has a median CTC of 819,999, which is slightly higher than Cluster 2 but lower than Cluster 1. - **Range:** The CTC ranges from 100,000 to 15,000,000, indicating a wide spread. This cluster also includes a mix of individuals with varying salary levels.

```
[132]: scaled_df.shape
```

```
[132]: (172443, 57)
```


0.1.17 Hierarchical Clustering

```
[135]: sampled_df = scaled_df.sample(55000)
sampled_df = sampled_df.drop(columns=['kM_cluster'])
```

```
[136]: # Hierarchical Clustering - Dendrogram

cluster_0 = sampled_df.values
Z = linkage(cluster_0, method='ward')

plt.figure(figsize=(10, 6))
dendrogram(Z)
plt.title('Dendrogram')
plt.xlabel('Clusters')
plt.ylabel('Euclidean Distance')
plt.show()
```



```
[144]: # Cut the dendrogram to form 3 clusters
num_clusters = 3
clusters = fcluster(Z, num_clusters, criterion='maxclust')

# Add the cluster labels to the subset DataFrame
sampled_df['Cluster'] = clusters
```

0.1.18 Actionable Insights and Recommendations

- **Targeted Interventions:**
 - For learners in the low category, consider providing additional training and resources to help them move to higher CTC categories.
 - For learners in the medium category, focus on career development and opportunities for promotions to help them achieve higher CTC values.
 - For learners in the high category, provide support for continued professional growth and retention strategies to maintain their high performance.
- **Role Transition:**
 - For learners experiencing the plateau effect, consider transitioning to higher-level roles (e.g., from software engineer to technical lead) to achieve further salary growth.
- **Negotiation Skills:**
 - Equip learners with negotiation skills to help them secure better salaries, especially when they have significant experience or specialized skills.